

Distributions

- Normal:** $X \sim N(\mu, \sigma^2)$, pdf: $f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$, mean: μ , variance: σ^2 . **SM:** Box-Muller Transform: If $U_1, U_2 \sim Uniform(0, 1)$ i.i.d., then $Z_1 = \sqrt{-2 \ln U_1} \cos(2\pi U_2)$ and $Z_2 = \sqrt{-2 \ln U_1} \sin(2\pi U_2)$ are i.i.d. $N(0, 1)$. To get $N(\mu, \sigma^2)$, use $X = \sigma Z + \mu$
- Bernoulli:** $X \sim Bern(p)$, pmf: $P(X = 1) = p$, $P(X = 0) = 1 - p$, mean: p , variance: $p(1 - p)$. **SM:** If $U \sim Uniform(0, 1)$, then $X = 1$ if $U \leq p$, else $X = 0$
- Binomial:** $X \sim Bin(n, p)$, pmf: $P(X = k) = \binom{n}{k} p^k (1-p)^{n-k}$, mean: np , variance: $np(1-p)$. Note: $Bin(n, p) = \sum_{i=1}^n Bern(p)$. **SM:** If $U_i \sim Uniform(0, 1)$ i.i.d., then $X = \sum_{i=1}^n I(U_i \leq p)$
- Multinomial** $\sim Mult(n, p_1, p_2, \dots, p_k)$, pmf: $P(X_1 = x_1, X_2 = x_2, \dots, X_k = x_k) = \frac{n!}{x_1!x_2!\dots x_k!} p_1^{x_1} p_2^{x_2} \dots p_k^{x_k}$, mean: $E[X_i] = np_i$, variance: $Var(X_i) = np_i(1-p_i)$, covariance: $Cov(X_i, X_j) = -np_i p_j$ for $i \neq j$. **SM:** If $U_i \sim Uniform(0, 1)$ i.i.d., then for each U_i , assign it to category j if $\sum_{m=1}^{j-1} p_m < U_i \leq \sum_{m=1}^j p_m$, then X_j is the count of assignments to category j
- Exponential:** $X \sim Exp(\lambda)$, pdf: $f(x) = \lambda e^{-\lambda x}$ for $x \geq 0$, mean: $\frac{1}{\lambda}$, variance: $\frac{1}{\lambda^2}$. **SM:** If $U \sim Uniform(0, 1)$, then $X = -\frac{1}{\lambda} \ln(U)$
- Poisson:** $X \sim Poisson(\lambda)$, pmf: $P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}$ for $k = 0, 1, 2, \dots$, mean: λ , variance: λ . Note: $Poisson(\lambda)$ with Exponential inter-arrival times. **SM:** $U \sim U(0, 1)$ set $f = e^{-\lambda}$, $k = 0$, $F = f$, while $F < u$ do $k = k + 1$, $f = f \frac{\lambda}{k}$, $F = F + f$, return k
- Chi-Squared:** $X \sim \chi_k^2$, pdf: $f(x) = \frac{1}{2^{k/2}\Gamma(k/2)} x^{(k/2)-1} e^{-x/2}$ for $x \geq 0$, mean: k , variance: $2k$. Note: If $Z_i \sim N(0, 1)$ i.i.d., then $\sum_{i=1}^k Z_i^2 \sim \chi_k^2$. **SM:** $z_k \sim N(0, 1)$ $x = (z_1 + \lambda)^2 + \sum_{i=2}^k z_i^2$
- t-Distribution:** $X \sim t_k$, pdf: $f(x) = \frac{\Gamma(\frac{k+1}{2})}{\sqrt{k\pi}\Gamma(\frac{k}{2})} (1 + \frac{x^2}{k})^{-\frac{k+1}{2}}$, mean: 0 for $k > 1$, variance: $\frac{k}{k-2}$ for $k > 2$. Note: If $Z \sim N(0, 1)$ and $V \sim \chi_k^2$ independent, then $\frac{Z}{\sqrt{V/k}} \sim t_k$. **SM:** Use the property of $T \sim N/\sqrt{(X/k)}$ where $N \sim N(0, 1)$ and $X \sim \chi_k^2$ independent. Generate N using Box-Muller and X using Chi-Squared SM.
- F-Distribution:** $X \sim F_{d_1, d_2}$, pdf: $f(x) = \frac{\sqrt{\frac{(d_1 x)^{d_1} d_2^{d_2}}{(d_1 x + d_2)^{d_1 + d_2}}}}{x B(\frac{d_1}{2}, \frac{d_2}{2})}$ for $x \geq 0$, mean: $\frac{d_2}{d_2 - 2}$ for $d_2 > 2$, variance: $\frac{2d_2^2(d_1 + d_2 - 2)}{d_1(d_2 - 2)^2(d_2 - 4)}$ for $d_2 > 4$. Note: If $U_1 \sim \chi_{d_1}^2$ and $U_2 \sim \chi_{d_2}^2$ independent, then $\frac{(U_1/d_1)}{(U_2/d_2)} \sim F_{d_1, d_2}$. **SM:** Generate U_1 and U_2 using Chi-Squared SM and use the property above.
- Gamma:** $X \sim Gamma(\alpha, \beta)$, pdf: $f(x) = \frac{1}{\beta^\alpha \Gamma(\alpha)} x^{\alpha-1} e^{-x/\beta}$ for $x \geq 0$, mean: $\alpha\beta$, variance: $\alpha\beta^2$. **SM:** If α is an integer, generate α i.i.d. $Exp(1/\beta)$ and sum them. If not integer, use acceptance-rejection or other methods.
- Beta:** $X \sim Beta(\alpha, \beta)$, pdf: $f(x) = \frac{x^{\alpha-1}(1-x)^{\beta-1}}{B(\alpha, \beta)}$ for $0 \leq x \leq 1$, mean: $\frac{\alpha}{\alpha+\beta}$, variance: $\frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$. **SM:** If $U_1 \sim Gamma(\alpha, 1)$ and $U_2 \sim Gamma(\beta, 1)$ independent, then $X = \frac{U_1}{U_1+U_2} \sim Beta(\alpha, \beta)$
- Inverse CDF:** For any $X \sim F(x)$, if $U \sim Uniform(0, 1)$, then $X = F^{-1}(U)$ has distribution $F(x)$. Use when F^{-1} is available in closed form.

Sampling Methods

- Rejection Sampling:** $X \sim f(x)$, $g(x)$ is proposal distribution with $f(x) \leq Mg(x)$ for all x , generate $Y \sim g(y)$ and $U \sim Uniform(0, 1)$, accept Y if $U \leq \frac{f(Y)}{Mg(Y)}$, else reject and repeat.
- Importance Sampling:** $X \sim f(x)$ where $f(x) = \frac{h(x)}{\int h(x)dx}$, $g(x)$ is proposal distribution, generate $X_i \sim g(x)$ i.i.d. for $i = 1, 2, \dots, n$. Then sample X from Y_i with weight $w_i = \frac{f(X_i)}{g(X_i)} = \frac{h(x)}{g(x)}$ normalized so that $\sum_{i=1}^n w_i = 1$. We can use this for estimating expectations: $E_f[t(X)] = \int t(x)f(x)dx = \int t(x)\frac{f(x)}{g(x)}g(x)dx = E_Z[t(Z)\frac{f(Z)}{g(Z)}]Z \sim g(x) \approx \frac{1}{n} \sum_{i=1}^n t(z_i)\frac{f(z_i)}{g(z_i)}$
- Gibbs Sampling:** To sample from joint distribution $f(x, y)$, initialize $X^{(0)}$ and $Y^{(0)}$, then for $i = 1, 2, \dots, n$, sample $X^{(i)} \sim f(x|Y^{(i-1)})$ and $Y^{(i)} \sim f(y|X^{(i)})$. The samples $(X^{(i)}, Y^{(i)})$ converge to the joint distribution $f(x, y)$ as $n \rightarrow \infty$. Need to do burn-in and thinning to reduce autocorrelation. For multi-dimensions, sample each variable in turn conditioned on the others with full conditional distributions.

MCMC & Metropolis-Hastings

- MC Method/Integration :** To estimate $I = \int_a^b f(x)dx$, generate $U_i \sim Uniform(a, b)$ i.i.d. for $i = 1, 2, \dots, n$, then $\hat{I} = \frac{b-a}{n} \sum_{i=1}^n f(U_i)$
- MCMC:** To sample from target distribution $f(x)$, construct a Markov chain with transition kernel $P(x, y)$ such that $f(x)$ is the stationary distribution. Run the chain for a long time and use the samples to estimate expectations. Converges requires π -irreducibility, aperiodic, and invariance distribution.
- Metropolis-Hastings:** To sample from target distribution $f(x)$, choose $g(\cdot|x)$. Init x_0 . For $i = 1, 2, \dots, n$, generate $Y \sim g(\cdot|x_{i-1})$, compute acceptance ratio $\alpha = \min(1, \frac{f(Y)g(x_{i-1}|Y)}{f(x_{i-1})g(Y|x_{i-1})})$, $x_i = Y$ with probability α , else $x_i = x_{i-1}$. The samples x_i converge to $f(x)$ as $n \rightarrow \infty$.
- Metropolis Algorithm:** Special case of Metropolis-Hastings where $g(y|x) = g(x|y)$ (symmetric proposal). Acceptance ratio simplifies to $\alpha = \min(1, \frac{f(Y)}{f(x_{i-1})})$.
- Independent MH:** Special case of Metropolis-Hastings where $g(y|x) = g(y)$ (independent proposal). Acceptance ratio is $\alpha = \min(1, \frac{f(Y)g(x_{i-1})}{f(x_{i-1})g(Y)})$.

Correlation Coefficient

Pearson Correlation Coefficient: For random variables X and Y , $\rho_p(X, Y) = \frac{\text{Cov}(X, Y)}{\sqrt{\text{Var}(X)\text{Var}(Y)}} = \frac{E[(X - E[X])(Y - E[Y])]}{\sqrt{E[(X - E[X])^2]E[(Y - E[Y])^2]}}$ Used for measuring linear relationship between variables. Sample: $\hat{\rho}_p = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2 \sum_{i=1}^n (y_i - \bar{y})^2}}$

Spearman's Rank Correlation Coefficient: For random variables X and Y , $\rho_s(X, Y) = \rho_p(F_X(z), F_Y(z)) = 12 \int \int F_X F_Y f(x, y) dx dy - 3$ where F_X and F_Y are the CDFs of X and Y . Measures monotonic relationship. Sample: $\hat{\rho}_s = \frac{\frac{1}{n} \sum r_i^{(x)} r_i^{(y)} - \bar{r}^{(x)} \bar{r}^{(y)}}{s_{r^{(x)}} s_{r^{(y)}}}$ Where $r_i^{(x)}$ is the rank of x_i among x_1, x_2, \dots, x_n

Kendall's Tau: For random variables X and Y , $\tau = P((X_1 - X_2)(Y_1 - Y_2) > 0) - P((X_1 - X_2)(Y_1 - Y_2) < 0)$ Used for non-linear relationships. Sample: for each pair (x_i, y_i) and (x_j, y_j) , count concordant pairs C and discordant pairs D , then $\hat{\tau} = \frac{C - D}{\binom{n}{2}} = 4 \int \int F(x, y) f(x, y) dx dy - 1$

Copulas

Copula: $C(t, s) : [0, 1]^2 \rightarrow [0, 1]$ to combine the marginal distributions $F_{X_i}(x)$ with correlations to form joint distribution $F_{X_1, X_2}(x_1, x_2) = C(F_{X_1}(x_1), F_{X_2}(x_2))$ Note that $C(t, s) = F(F^{-1}(t), F^{-1}(s)) \iff F(x, y) = C(F(x), F(y))$ **Sklar's Theorem:** For any multivariate distribution F with marginals F_1, F_2, \dots, F_n , there exists a copula C such that $F(x_1, x_2, \dots, x_n) = C(F_1(x_1), F_2(x_2), \dots, F_n(x_n))$. If F_i are continuous, then C is unique. Conversely, for any copula C and marginals F_i , the function F defined above is a multivariate distribution with marginals F_i .

Cholesky Factorization: For a positive definite matrix Σ , there exists a unique lower triangular matrix L such that $\Sigma = LL^T$. Can generate correlated normals: $\Sigma = [\sigma_{ij}]$ then $a_{ij} = \frac{\sigma_{ij} - \sum_{k=1}^{j-1} a_{ik}a_{jk}}{a_{jj}}$ for $i \geq j$, $a_{ii} = \sqrt{\sigma_{ii} - \sum_{k=1}^{i-1} a_{ik}^2}$ Algo: For $j = 1$ to n , for $i = j$ to n , compute $v_i = \sigma_{ij} - \sum_{k=1}^{j-1} a_{ik}a_{jk}$, then set $a_{jj} = \sqrt{v_j}$, and for $i = j + 1$ to n , set $a_{ij} = \frac{v_i}{a_{jj}}$.

Multivar Normal Sim: $X \sim N(\mu, \Sigma)$, $LL^T = \Sigma$ (Cholesky), $Z \sim N(0, I)$, then $X = LZ + \mu$

Gaussian Copula: Given correlation matrix Σ , the Gaussian copula is defined as $C(u_1, u_2, \dots, u_n) = \Phi_\Sigma(\Phi^{-1}(u_1), \Phi^{-1}(u_2), \dots, \Phi^{-1}(u_n))$ where Φ_Σ is the CDF of multivariate normal with covariance Σ and Φ^{-1} is the inverse CDF of standard normal.

Gaussian Copula Simulation: $Z_n \sim N(0, 1)$ iid $Z^* = LZ$, for i, n $u_i = \Phi(z_i^*/\sigma_i)$, $x_i = F_i^{-1}(u_i)$ This gives correlated samples x_i with marginals F_i and correlation structure from Σ .

Student t-Copula: Same as gaussian but add $\chi^2(\nu)$ scaling in $F(Z_i^*/(\sigma_i \sqrt{W/\nu}))$ where $W \sim \chi^2(\nu)$ independent.

Archimedian Copula: Given a continuous, strictly decreasing function $\phi : [0, 1] \rightarrow [0, \infty]$ with $\phi(1) = 0$, the Archimedian copula is defined as $C(u_1, u_2, \dots, u_n) = \phi^{-1}(\phi(u_1) + \phi(u_2) + \dots + \phi(u_n))$ where ϕ^{-1} is the pseudo-inverse of ϕ . Common examples include Clayton, Gumbel, and Frank copulas.

Bootstrap

Bootstrap: Using given data resample with replacement to create new datasets and estimate statistics. Algo: Take the bootstrap samples, calculate the statistic, order the statistics, and form confidence intervals.

Symmetric CI: $[\hat{\theta}_L, \hat{\theta}_U]$ where $L = \frac{\alpha}{2}N$, $U = (1 - \frac{\alpha}{2})N$

Asymmetric CI: $[2\hat{\theta} - \hat{\theta}_{(U)}, 2\hat{\theta} - \hat{\theta}_{(L)}]$.

Bootstrap CLT $\hat{\theta}^* - \hat{\theta} \xrightarrow{d} \hat{\theta} - \theta_0 | \theta_0$ as $n \rightarrow \infty$

Bootstrap Residuals: Fit reg line, calculate residuals, resample residuals with replacement, create new response variable, refit reg line, repeat B times to get bootstrap estimates.

Bayes

Bayes Theorem: For events A and B with $P(B) > 0$, $P(A|B) = \frac{P(B|A)P(A)}{P(B)}$. For continuous random variables, $f_{X|Y}(x|y) = \frac{f_{Y|X}(y|x)f_X(x)}{f_Y(y)}$.

Prior : Initial belief about parameter θ before observing data, denoted as $p(\theta)$. **Conjugate**: post same family as prior, **Elicited**: from expert knowledge, **Non-informative**: vague prior Jefferys prior invariant under transformation.

Likelihood : Probability of observed data given parameter θ , denoted as $p(D|\theta)$. For i.i.d. data, $p(D|\theta) = \prod_{i=1}^n p(x_i|\theta)$.

Posterior : Updated belief about parameter θ after observing data, denoted as $p(\theta|D)$. Given by Bayes theorem: $p(\theta|D) = \frac{p(D|\theta)p(\theta)}{p(D)}$ where $p(D) = \int p(D|\theta)p(\theta)d\theta$.

Bayesian Inference: Use posterior distribution to make inferences about θ . Common summaries include posterior mean, median, mode, and credible intervals.

Bayes Factor $BF = \frac{P(D|M_1)}{P(D|M_2)}$ where $P(D|M_i) = \int P(D|\theta_i, M_i)P(\theta_i|M_i)d\theta_i$ is the marginal likelihood under model M_i . Used for model comparison. 1/10, 1/3, 1, 3, 10 str M2, mod, weak, weak, mod, str M1 evidence.

Bayes Example: Data 13/16 pref A to B. Prior $Beta(.5, .5)$, Likelihood $Bin(16, \theta)$, Posterior $Beta(13 + .5, 3 + .5) = Beta(13.5, 3.5)$.

$BF = \frac{\int_z^1 f(y|\theta)\pi(\theta)d\theta}{\int_0^z f(y|\theta)\pi(\theta)d\theta}$ where z = critical value for H0: $\theta \geq z$ vs H1: $\theta < z$

Linear Model Bayes: $Y = X\beta + \epsilon$, $\epsilon \sim N(0, \sigma^2 I)$, observations $y|\beta, \sigma^2 \sim N(X\beta, \sigma^2 I)$, Prior: $\beta|\sigma^2 \sim N(\beta_0, \sigma^2 B_0)$, $\sigma^2 \sim \mathcal{G}(c, C_0)$, Posterior: $p(\beta, \sigma^2|y) \propto p(y|\beta, \sigma^2)p(\beta|\sigma^2)p(\sigma^2)$